



Generative Artificial Intelligence

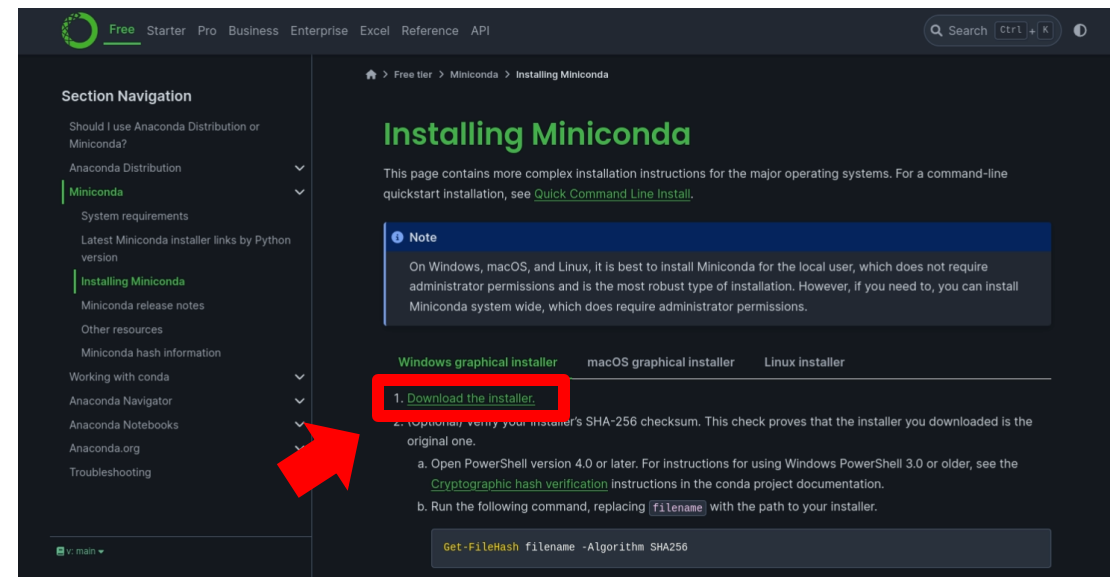
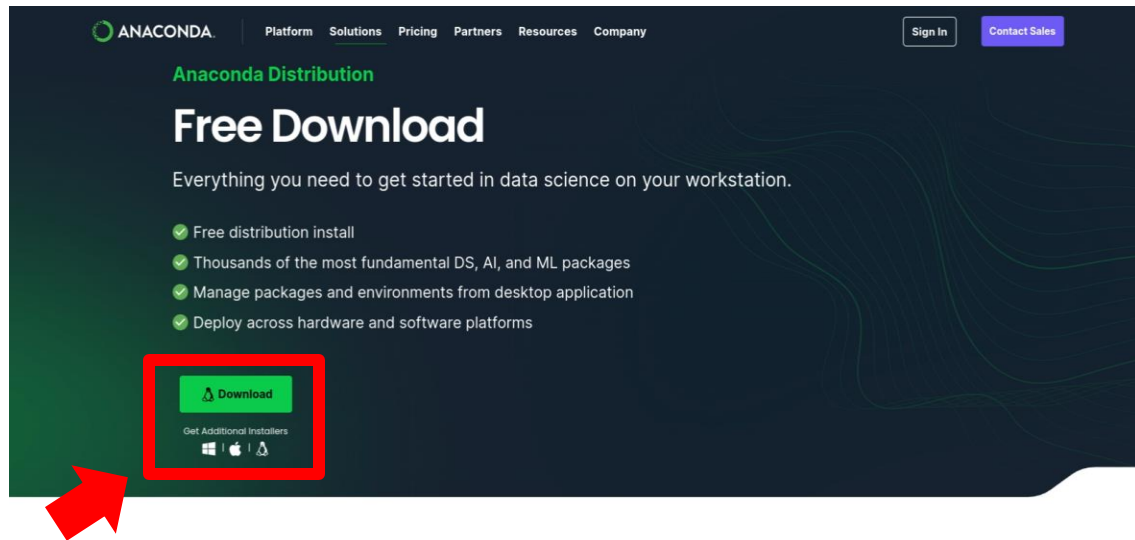
Pytorch tutorial



Environment setup

下載 anaconda (<https://www.anaconda.com/download>)


或者 miniconda (<https://docs.anaconda.com/free/miniconda/miniconda-install/>)



Environment setup

開啟Anaconda prompt

當前環境



Basic Commands

安裝完anaconda以後，開啟anaconda prompt，可以通過commands操作環境。
也可以直接通過圖形化介面navigator操作

1. 創建新的python環境：
`conda create -n (name) python=3.X.X`
2. 激活python環境
`conda activate (name)`
3. 退出環境
`conda deactivate`
4. 查看所有環境
`conda env list`



Install Pytorch

安裝pytorch: <https://pytorch.org/>

PyTorch Get Started Ecosystem ▾ Edge ▾ Blog Tutorials Docs ▾ Resources ▾ GitHub 🔍

NOTE: Latest PyTorch requires Python 3.8 or later. For more details, see Python section below.

PyTorch Build

Your OS

Package

Language

Compute Platform

Run this Command:

Previous versions of PyTorch >

Stable (2.2.1) Preview (Nightly)

Linux Mac Windows

Conda Pip LibTorch Source

Python C++ / Java

CUDA 11.8 CUDA 12.1 ROCm 5.7 CPU

```
conda install pytorch torchvision torchaudio pytorch-cuda=11.8 -c pytorch -c nvidia
```

aws Amazon Web Services >

Google Cloud Platform >

Microsoft Azure >

Install Pytorch

也可以選擇更早版本的，主要是cuda的版本電腦的硬體要能夠支持

The screenshot shows the PyTorch website's installation guide. The navigation bar includes links for Get Started, Ecosystem, Edge, Blog, Tutorials, Docs, Resources, and GitHub. A note states: "NOTE: Latest PyTorch requires Python 3.8 or later. For more details, see Python section below." The main content area features a table of installation options:

PyTorch Build	Stable (2.2.1)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.8	CUDA 12.1	ROCm 5.7	CPU

Below the table, the command to run is: `conda install pytorch torchvision torchaudio pytorch-cuda=11.8 -c pytorch -c nvidia`

On the right side, there are links to Amazon Web Services, Google Cloud Platform, and Microsoft Azure.

At the bottom left, a red box highlights the link "Previous versions of PyTorch >", with a red arrow pointing to it.

Training

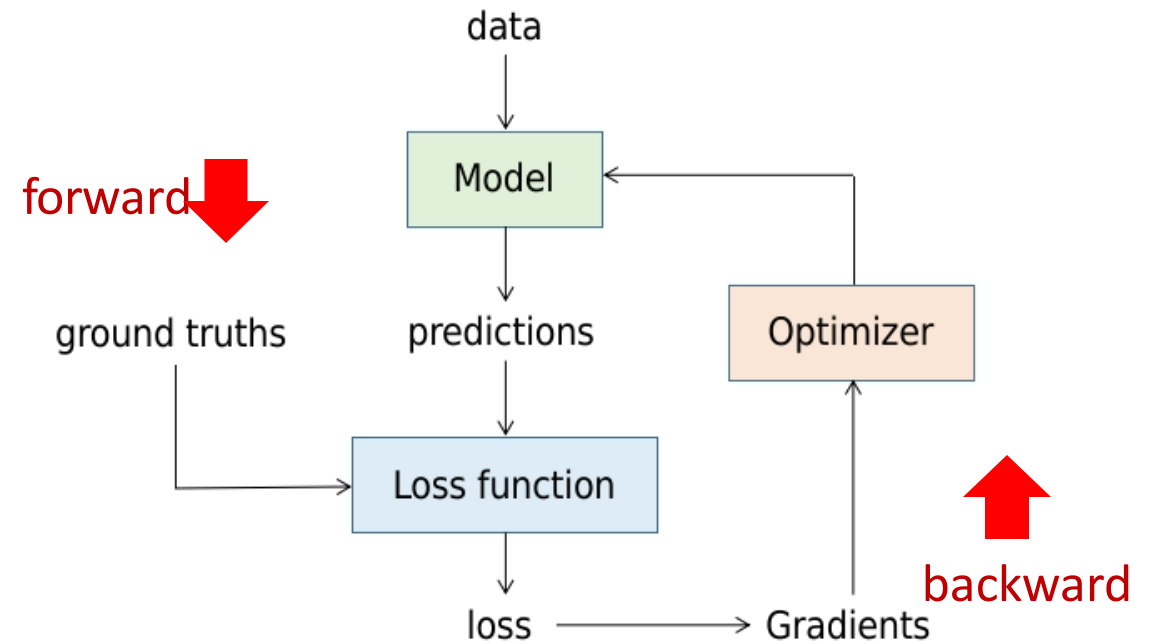
Forward: Obtain output logits

`model.forward()`

Backward: Back-propagation

Use Autograd

`loss.backward()`

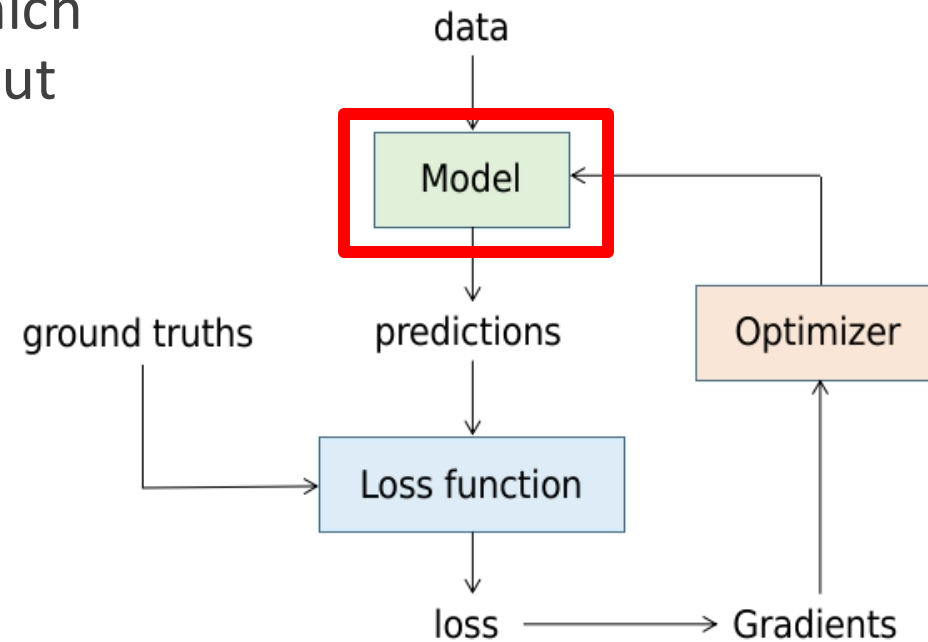
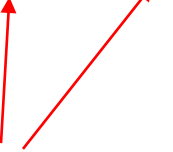


What is a Model

A model is like a special function which define the relationship between input and output.

```
def func(x):  
    return 3*x+4
```

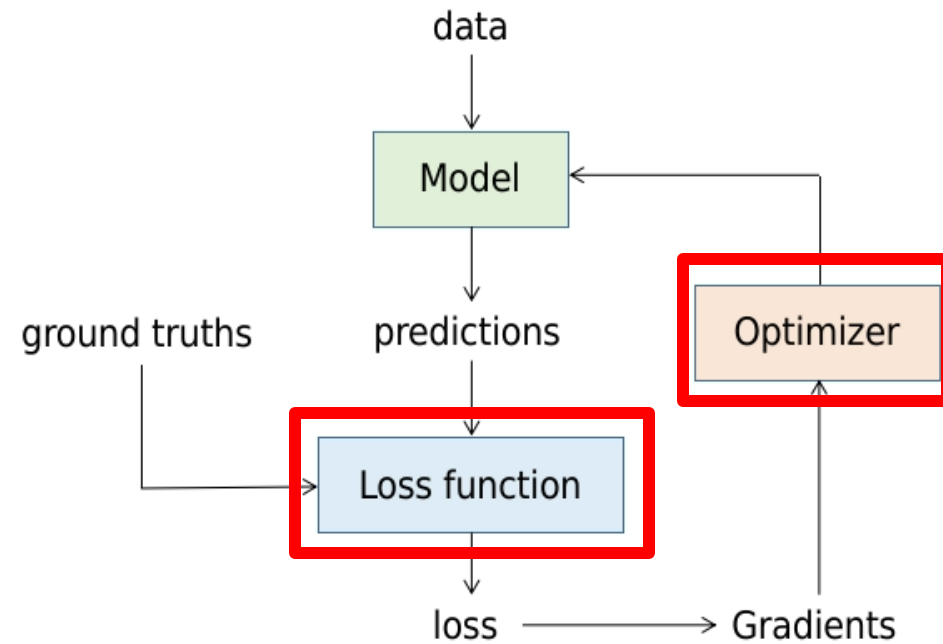
parameters



What is Loss function&Optimizer

Loss Function defines the distance between model's prediction & expected answer (ground truth).

Optimizer defines the way to optimize the parameters.



A simple example

We have a linear "model": $y = 2x + 1$

Input a data pair: $(1, 1)$

When $x=1$ we get: $y = 2*1 + 1 = 3$

Loss function: What is the difference between expected value 1 and model output 3? $\text{loss} = 3 - 1 = 2$?

Optimizer: How to fix the difference between expected value 1 and model output 3?



Forward & Back-propagation

`model.forward()`


$l1 = \text{input} \times w1$

$l2 = l1 + w2$

$l3 = l1 \times w3$

$l4 = l2 \times l3$

$\text{loss} = \text{mean}(l4)$



Forward & Back-propagation

`loss.backward()`

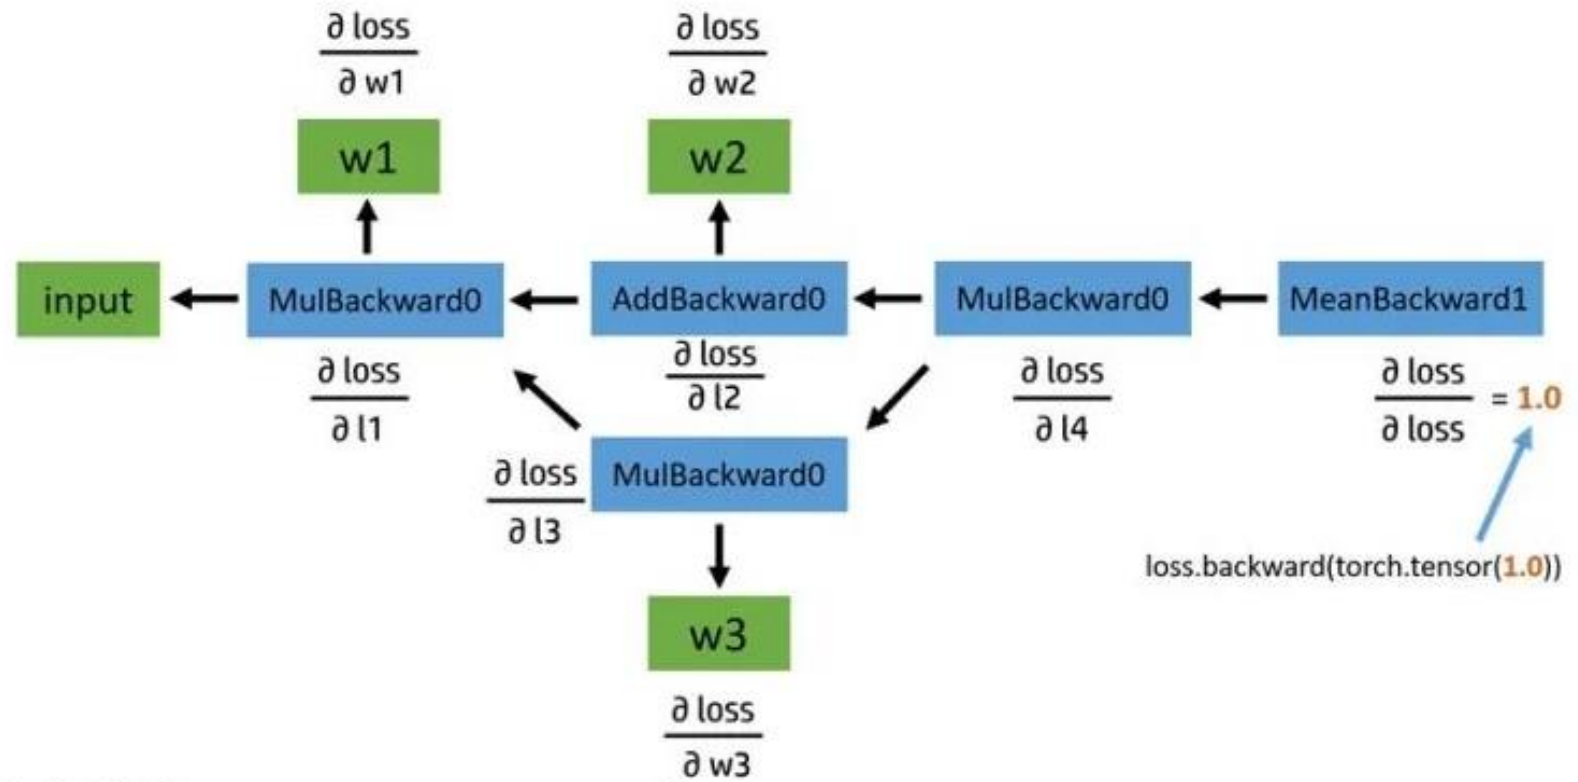
`l1 = input x w1`

`l2 = l1 + w2`

`l3 = l1 x w3`

`l4 = l2 x l3`

`loss = mean(l4)`



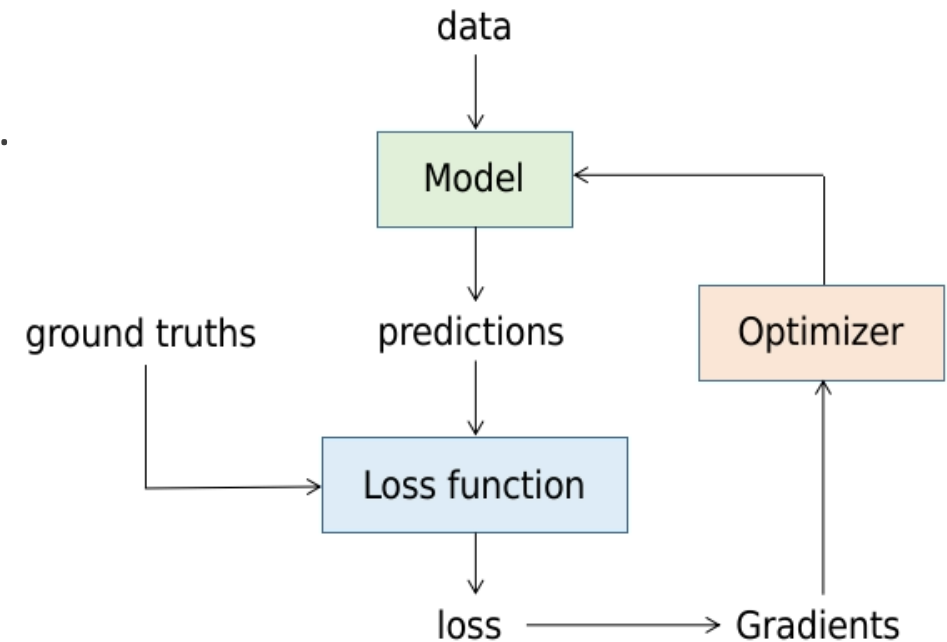
Problem description

In a deep learning project, we need to:

1. Define a model
 - Input a batch of data, and compute the results.
2. Train the model
 - Record the derivation procedures.
 - Back propagate & Compute the gradients.
 - Optimize the parameters.
 - GPU acceleration.


➡ A special data structure was invented.

Called "**Tensor**"(張量)



Tensors

- Tensors are very similar to arrays and matrices.
- Tensors can run on GPUs or other hardware accelerators.



Initialize a Tensor

There are many ways to initialize a pytorch tensor:

From existed python list

```
T = torch.tensor([[1, 2], [3, 4]])
```

From numpy

```
T = torch.from_numpy(numpy.array([[1, 2], [3, 4]]))
```

Randomly initialized

```
T = torch.randn(size=(2,2))
```

Initialize a Tensor

Other methods

All of the elements are 1

```
T = torch.ones(size=(2,2))
```

All of the elements are 0

```
T = torch.zeros(size=(2,2))
```

Diagonalized matrix

```
T = torch.eye(n=2, m=3)
```



Tensors

Initialize: `T = torch.tensor([[1, 2], [3, 4]])`

- Data (Scalar only)

`t = torch.tensor(1.1), t.item() => 1.1`

- Shape (return a tuple)

`T.shape => torch.Size([2, 2])`

- Data type

`T.dtype => torch.int64`

Tensors

- Device
`T.device => device(type='cpu')`
- Require gradient
`T.requires_grad => False`
- Gradient
`T.grad => tensor([...])`

Matrix operations

Tensor & scalar

```
A = torch.tensor([[1, 2], [3, 4]])
```

```
A+1 => tensor([[2, 3], [4, 5]])
```

```
A*2 => tensor([[2, 4], [6, 8]])
```

符號	意義
<code>torch.Tensor + scalar</code>	張量中的每個數值加上 <code>scalar</code>
<code>torch.Tensor - scalar</code>	張量中的每個數值減去 <code>scalar</code>
<code>torch.Tensor * scalar</code>	張量中的每個數值乘上 <code>scalar</code>
<code>torch.Tensor / scalar</code>	張量中的每個數值除以 <code>scalar</code>
<code>torch.Tensor // scalar</code>	張量中的每個數值除以 <code>scalar</code> 所得之商
<code>torch.Tensor % scalar</code>	張量中的每個數值除以 <code>scalar</code> 所得之餘數
<code>torch.Tensor ** scalar</code>	張量中的每個數值取 <code>scalar</code> 次方

The operation is applied equally to every elements in the tensor.



Matrix operations

```
A = torch.tensor([[1, 2], [3, 4]])
```

```
B = torch.tensor([[0, 1], [2, 3]])
```

+, -

```
A+B => tensor([[1, 3], [5, 7]])
```

```
A-B => tensor([[1, 1], [1, 1]])
```

Matrix multiply

```
A@B => tensor([[4, 7], [8, 15]])
```

Element-wise multiply

```
A*B => tensor([[0, 2], [6, 12]])
```

符號	意義
A + B	張量 A 中的每個數值加上張量 B 中相同位置的數值
A - B	張量 A 中的每個數值減去張量 B 中相同位置的數值
A * B	張量 A 中的每個數值乘上張量 B 中相同位置的數值
A / B	張量 A 中的每個數值除以張量 B 中相同位置的數值
A // B	張量 A 中的每個數值除以張量 B 中相同位置的數值所得之商
A % B	張量 A 中的每個數值除以張量 B 中相同位置的數值所得之餘數
A ** B	張量 A 中的每個數值取張量 B 中相同位置的數值之次方

Matrix operations

Concatenation

`torch.concat((A,B), dim=0) => tensor([[1, 2], [3, 4], [0, 1], [2, 3]])`

`torch.concat((A,B), dim=1) => tensor([[1, 2, 0, 1], [3, 4, 2, 3]])`

Squeeze, Unsqueeze

`A.unsqueeze(dim=-1) => tensor([[[1], [2]], [[3], [4]]])`

`A.squeeze(dim=-1) => tensor([[1, 2], [3, 4]])`

Transpose

`torch.transpose(A, dim0=0, dim1=1) => tensor([[1, 3], [2, 4]])`



Matrix operations

Reshape

A.view(4, 1) => tensor([[1], [2], [3], [4]])

A.view(1, 4) => tensor([[1, 2, 3, 4]])

A.reshape(4, 1) => tensor([[1], [2], [3], [4]])

A.reshape(1, 4) => tensor([[1, 2, 3, 4]])

Matrix operations

Note: All PyTorch computations are performed in a matrix operation mode.

e.g. Randomly initialize a tensor(100X100) and normalize it:

Sequentially: 0.127 s

Matrix Operation: 0.000088s

```
def sequential(mat : torch.tensor):  
    t = time.time()  
    average = torch.mean(mat)  
    standard = torch.std(mat)  
    output = torch.zeros([mat.shape])  
    for i in range(mat.shape[0]):  
        for j in range(mat.shape[1]):  
            output[i, j] = (mat[i, j]-average)/standard  
    return output, time.time() - t  
  
def parallel(mat : torch.tensor):  
    t = time.time()  
    average = torch.mean(mat)  
    standard = torch.std(mat)  
    output = (mat - average)/standard  
    return output, time.time() - t
```

Train a Model

Step1: Prepare the dataset

Step2: Construct the model

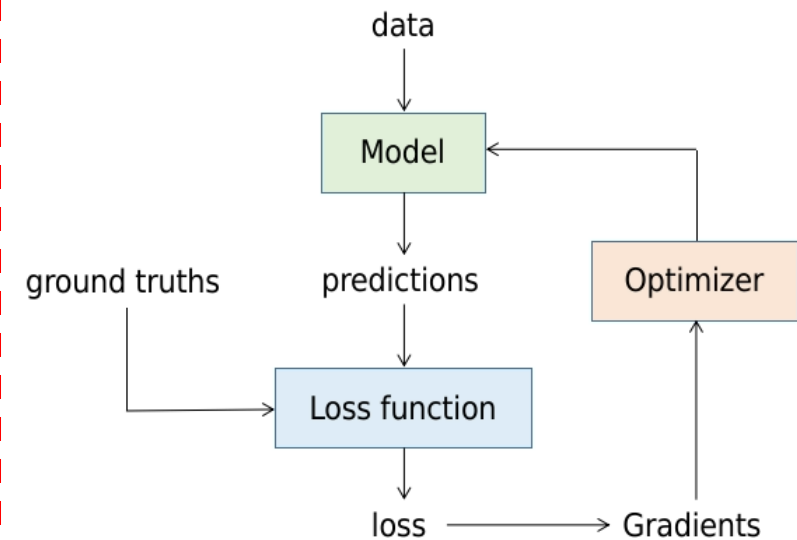
Step3: Define Optimizer

Step4: Define loss function

Step5: Train the model

Step6: Evaluate the model

1. clear gradients
2. Input data to the model
3. compute loss
4. compute gradients
5. optimize parameters
6. back to 1.



Training

1. clear gradients

`optimizer.zero_grad()`

2. Input data to the model

`output = model(**batch)`

3. compute loss

`loss = loss_fn(output)`

4. compute gradients

`loss.backward()`

5. optimize parameters

`optimizer.step()`


6. back to 1.



An Example of Linear Regression

Data distribution:

Use a line to represent these data



An Example of Linear Regression

Model → `# Linear regression model`
`model = nn.Linear(input_size, output_size)`

Loss → `# Loss and optimizer`
`criterion = nn.MSELoss()`

Optimizer → `optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)`

load data → `# Train the model`
`for epoch in range(num_epochs):`
 `# Convert numpy arrays to torch tensors`
 `inputs = torch.from_numpy(x_train)`
 `targets = torch.from_numpy(y_train)`

forward → `# Forward pass`
`outputs = model(inputs)`
`loss = criterion(outputs, targets)`

backward → `# Backward and optimize`
`optimizer.zero_grad()`
`loss.backward()`

optimize → `optimizer.step()`

`if (epoch+1) % 5 == 0:`
 `print ('Epoch [{}/{}], Loss: {:.4f}'.format(epoch+1, num_epochs, loss.item()))`



Check the parameters&gradients

```
inputs = torch.from_numpy(x_train)
targets = torch.from_numpy(y_train)

# Forward pass
outputs = model(inputs)
loss = criterion(outputs, targets)
```

```
optimizer.zero_grad()
print_grads(model)
```

```
weight: Parameter containing:
tensor([[0.4165]], requires_grad=True)
weight grad: None
```

```
bias: Parameter containing:
tensor([0.4819], requires_grad=True)
bias grad: None
```

Check the gradients

```
loss.backward()  
print_grads(model)
```

```
weight: Parameter containing:  
tensor([[0.4165]], requires_grad=True)  
weight grad: tensor([[10.0239]])
```

```
bias: Parameter containing:  
tensor([0.4819], requires_grad=True)  
bias grad: tensor([1.3666])
```

```
optimizer.step()  
print_grads(model)
```

```
weight: Parameter containing:  
tensor([[0.4065]], requires_grad=True)  
weight grad: tensor([[10.0239]])
```

```
bias: Parameter containing:  
tensor([0.4805], requires_grad=True)  
bias grad: tensor([1.3666])
```

```
optimizer.zero_grad()  
print_grads(model)
```

```
weight: Parameter containing:  
tensor([[0.4065]], requires_grad=True)  
weight grad: None
```


```
bias: Parameter containing:  
tensor([0.4805], requires_grad=True)  
bias grad: None
```



An Example of Linear Regression

Outputs:

```
Epoch [5/60], Loss: 11.2489
Epoch [10/60], Loss: 4.6657
Epoch [15/60], Loss: 1.9987
Epoch [20/60], Loss: 0.9182
Epoch [25/60], Loss: 0.4805
Epoch [30/60], Loss: 0.3031
Epoch [35/60], Loss: 0.2313
Epoch [40/60], Loss: 0.2021
Epoch [45/60], Loss: 0.1903
Epoch [50/60], Loss: 0.1855
Epoch [55/60], Loss: 0.1835
Epoch [60/60], Loss: 0.1827
```





NLP Tasks

In NLP tasks ...

`torch.utils.data.Dataset`




Text dataset
usually very big

`torch.utils..data.DataLoader`



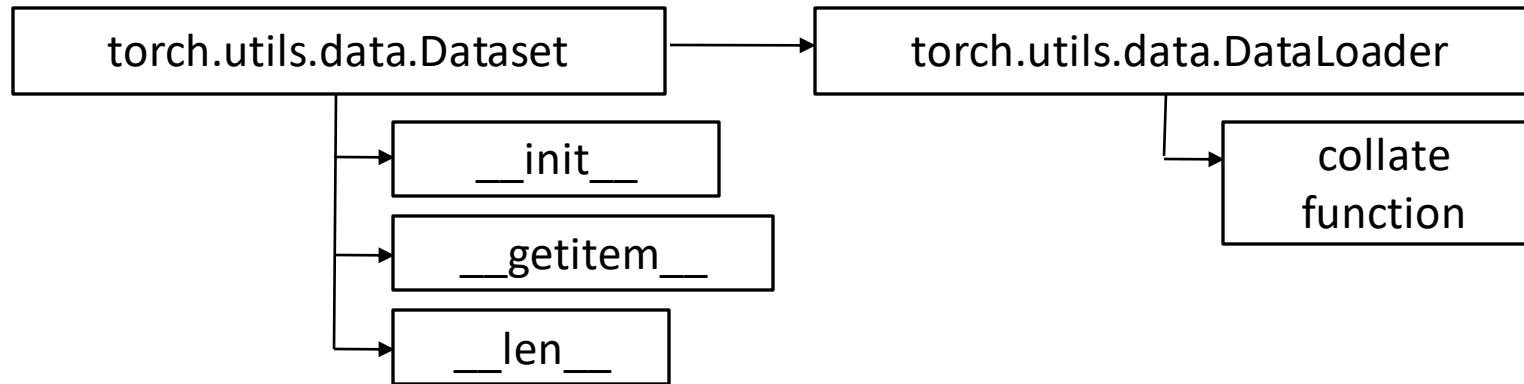
We need to divide the
dataset into numerous
data batches and gradually
feed them to a big
complex model.

`torch.nn.Module`



a big complex
language model.

Dataset & DataLoader



```
class myDataset(Dataset):
    def __init__(self, split):
        Super().__init__()
        self.data = ...

    def __getitem__(self, index):
        return self.data[index]

    def __len__(self):
        return len(self.data)
```

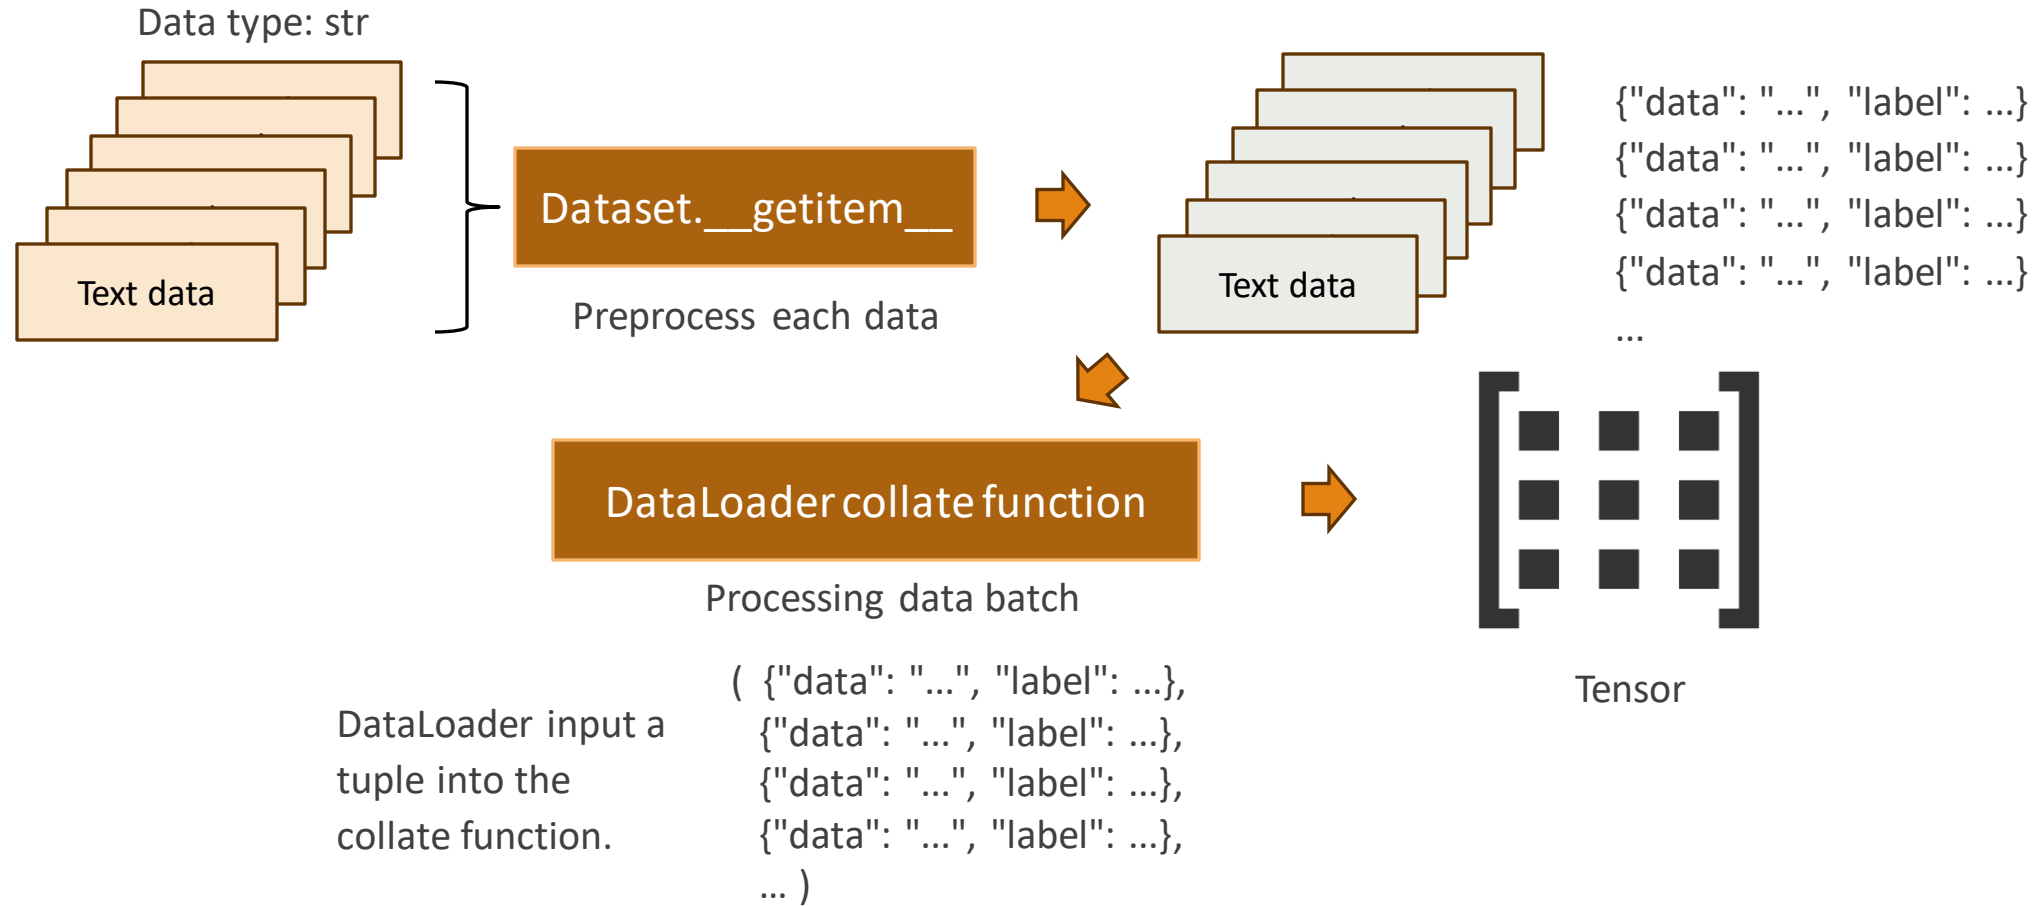
```
def get_batch(sample):
    ...

dl = DataLoader(myDataset(split=..., ), batch_size=...,
               collate_fn=get_batch, shuffle=...)

for batch in dl:
    ...
```

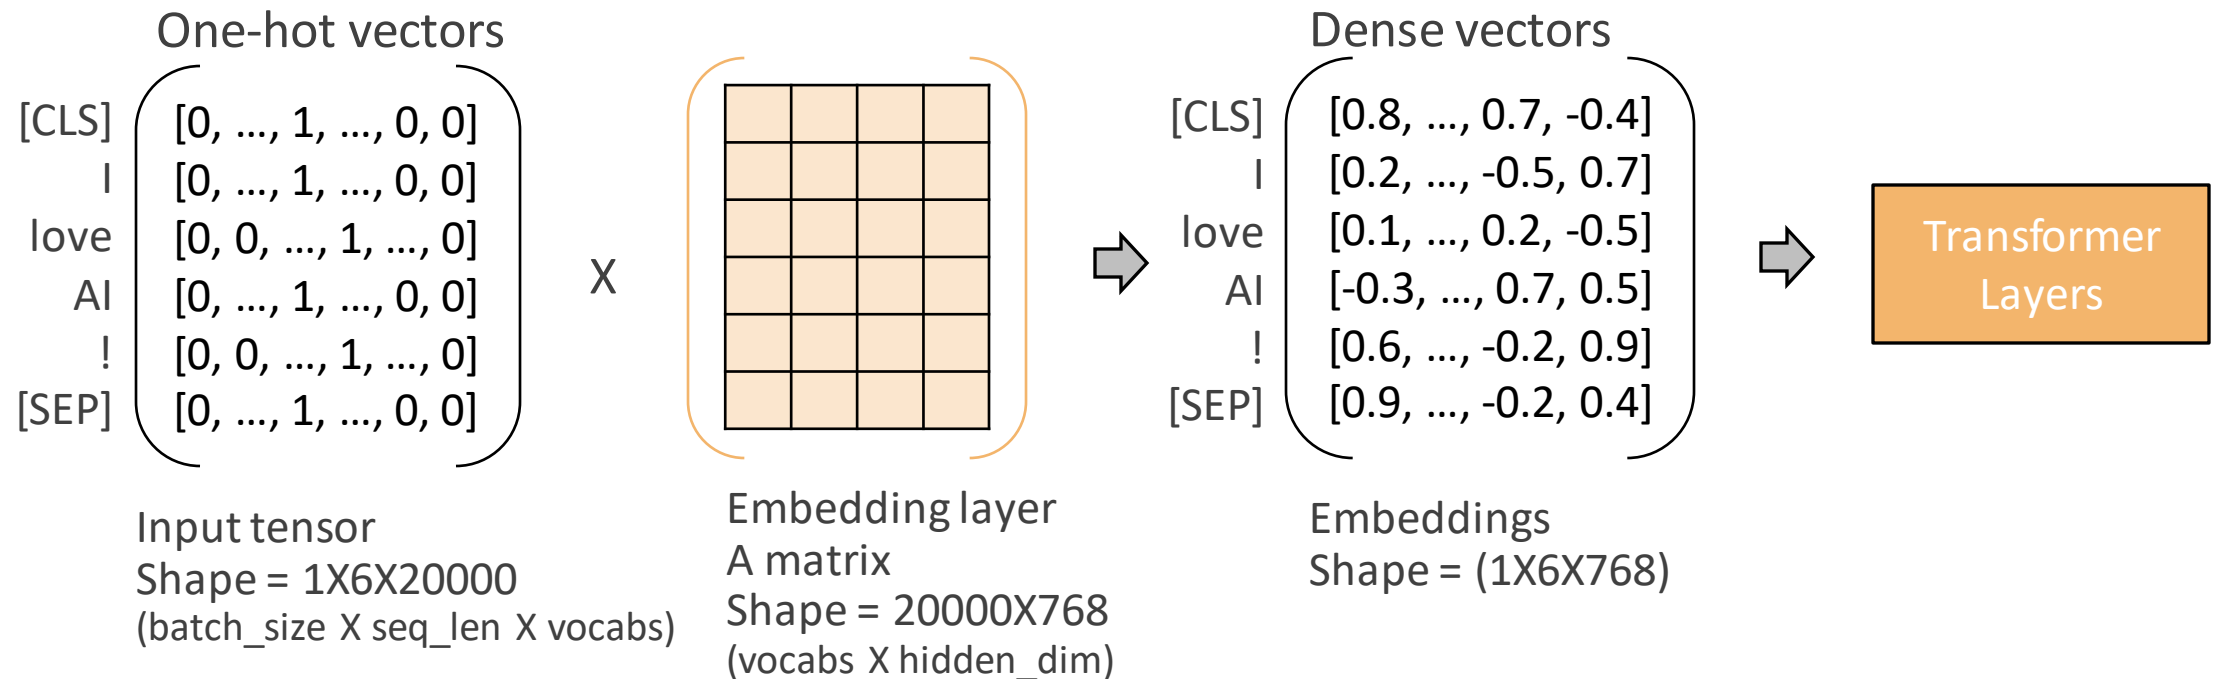


Data Flow in NLP tasks



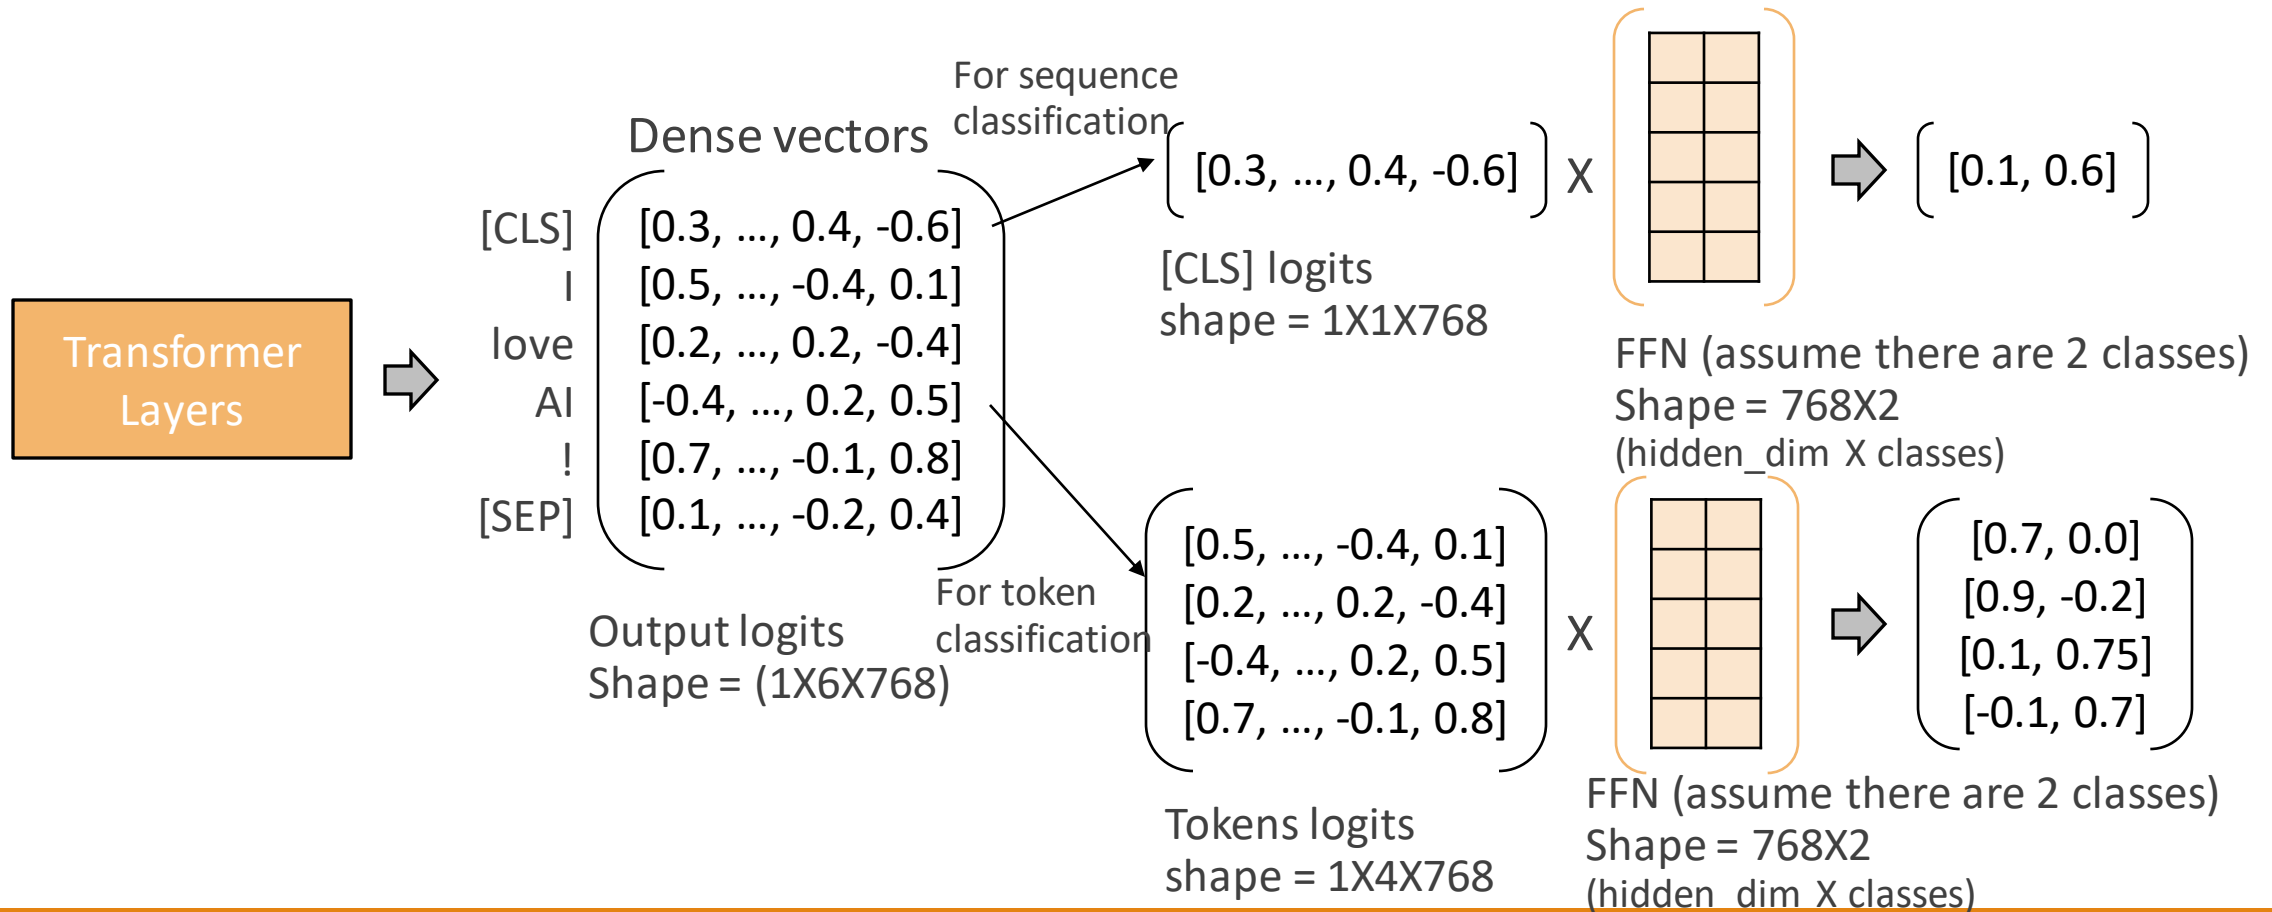
Data Flow in NLP tasks

We assume that `batch_size = 1`, `hidden_dim=768` and there are 20000 words in the dictionary

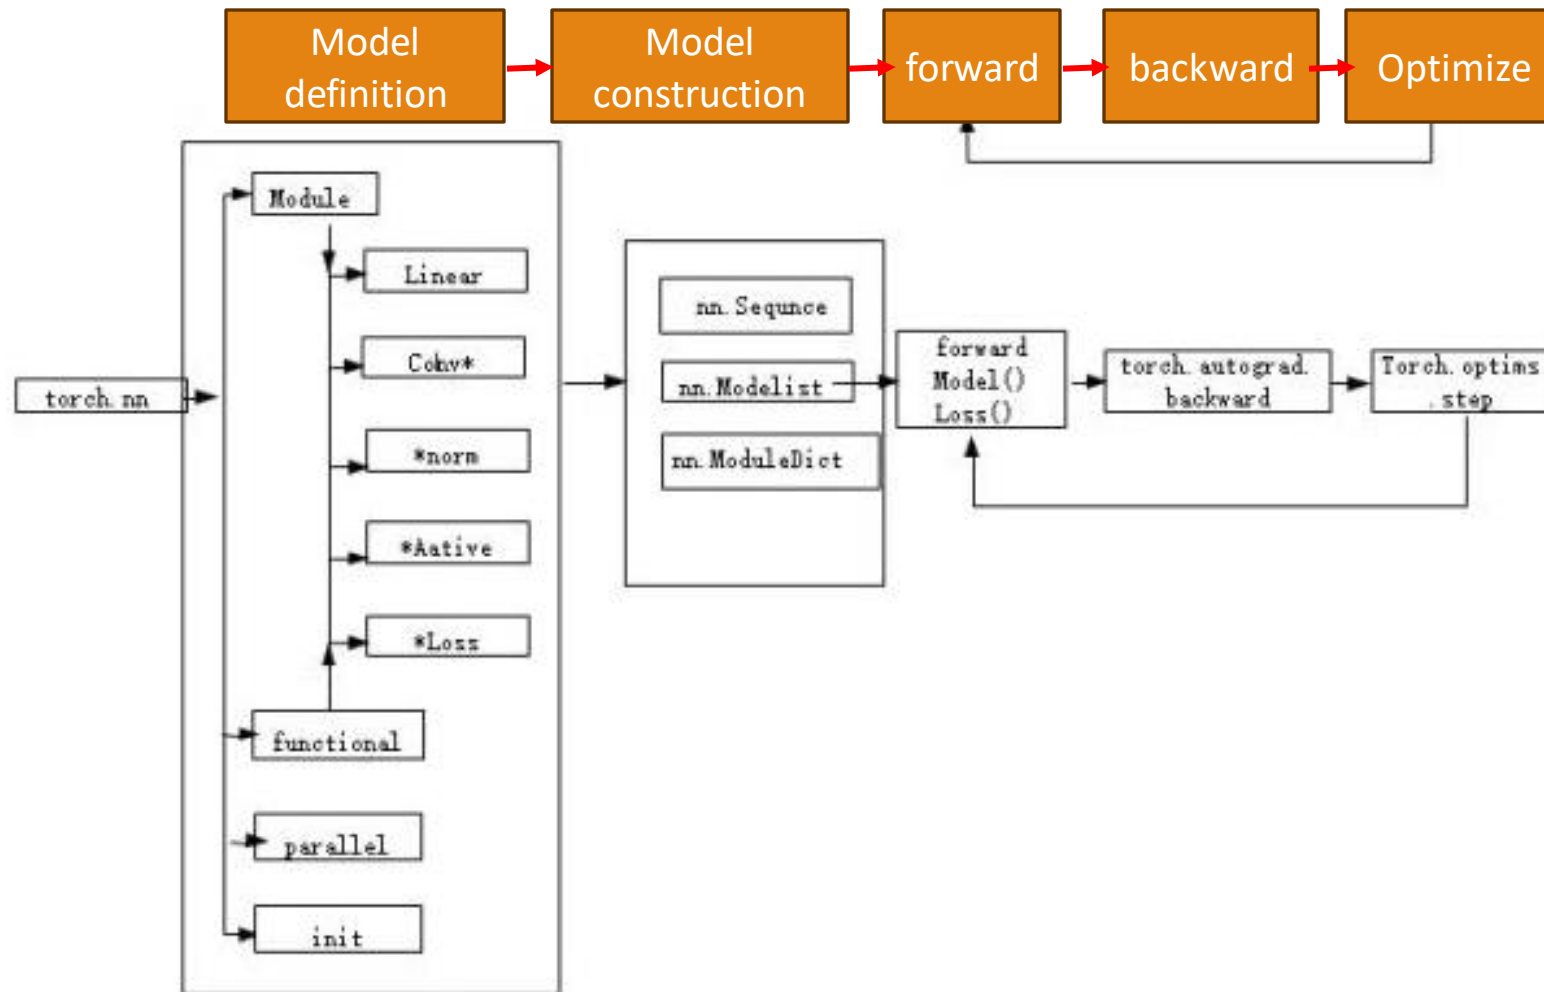


Data Flow in NLP tasks

We assume that `batch_size = 1`, `hidden_dim=768` and there are 20000 words in the dictionary



Construct & Train a Model



Construct a Model


A model is defined as:

```
class myModel(torch.nn.Module):  
    def __init__(self):  
        Super().__init__()  
        self.layers = ...  
  
    def forward(self, inputs):  
        ...  
        return ...
```

Obtain the output logits

```
model = myModel()  
logits = model(inputs)
```


`torch.nn.Module.__call__`



Pretrain & Finetuning

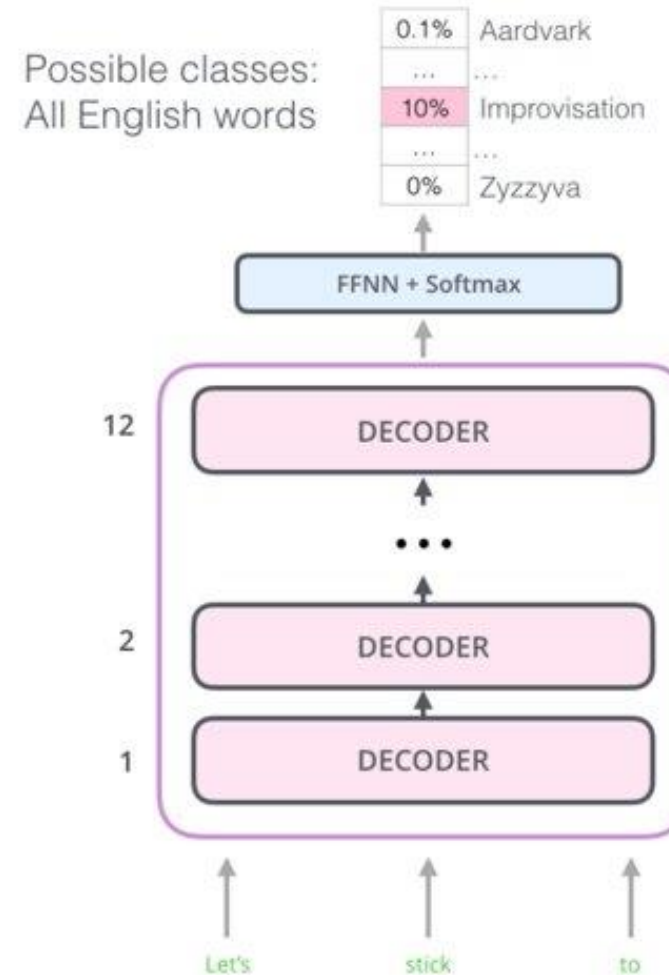

Prof said that transformers have to learn far more features than previous networks, so we have to pretrain the model.

Usually we do not have enough data to train the model directly to the target.

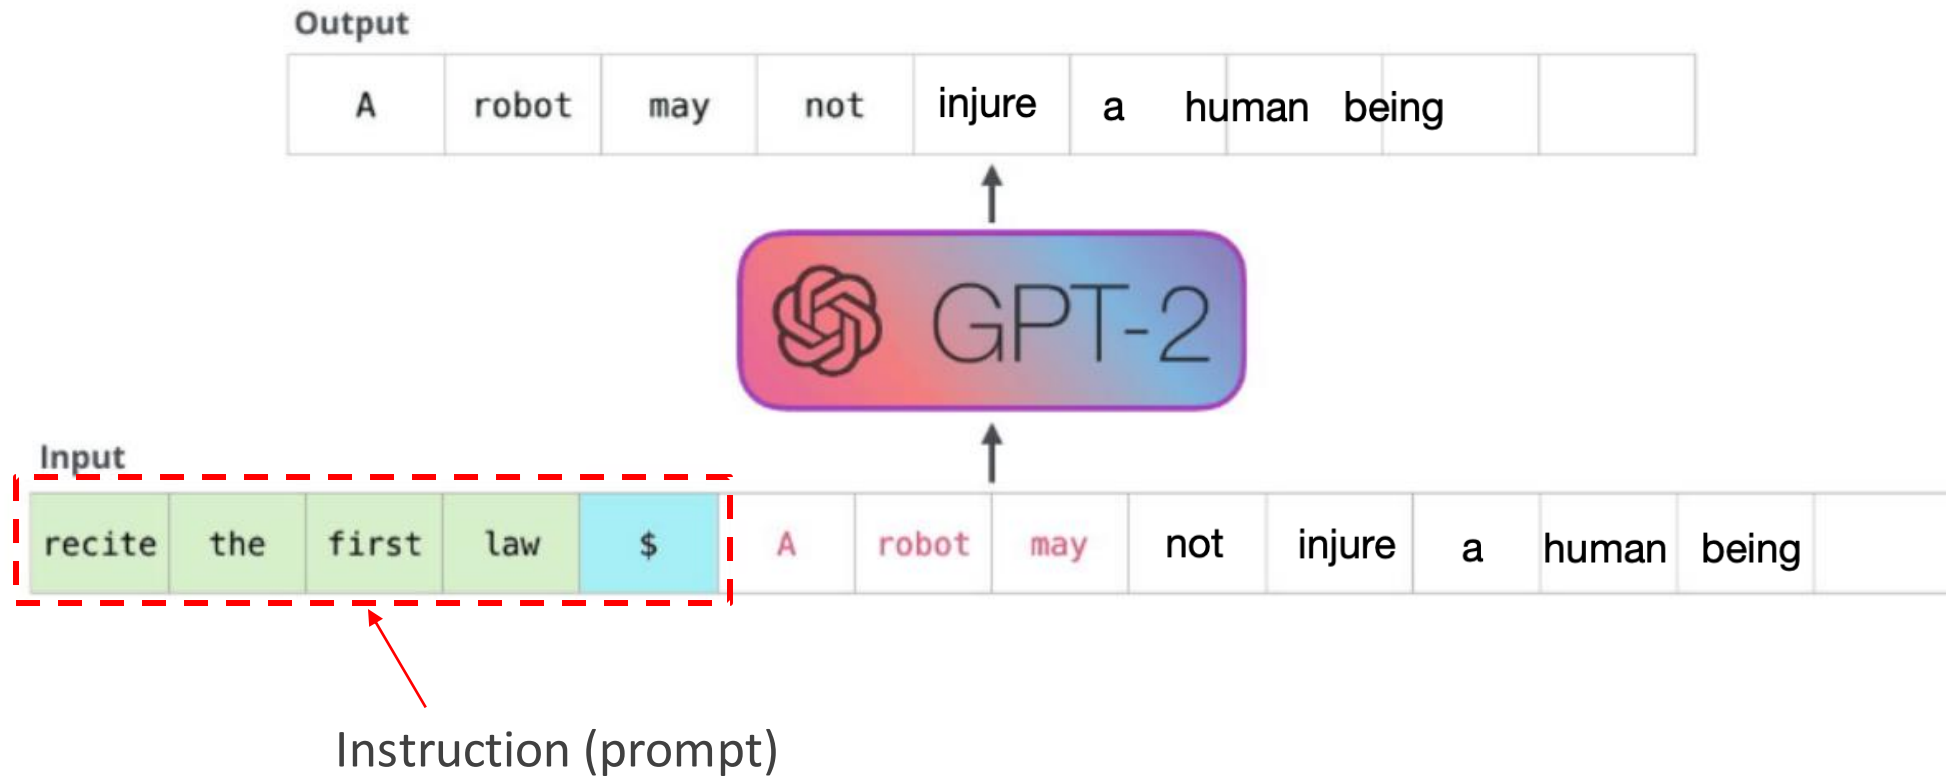


GPT2 Review

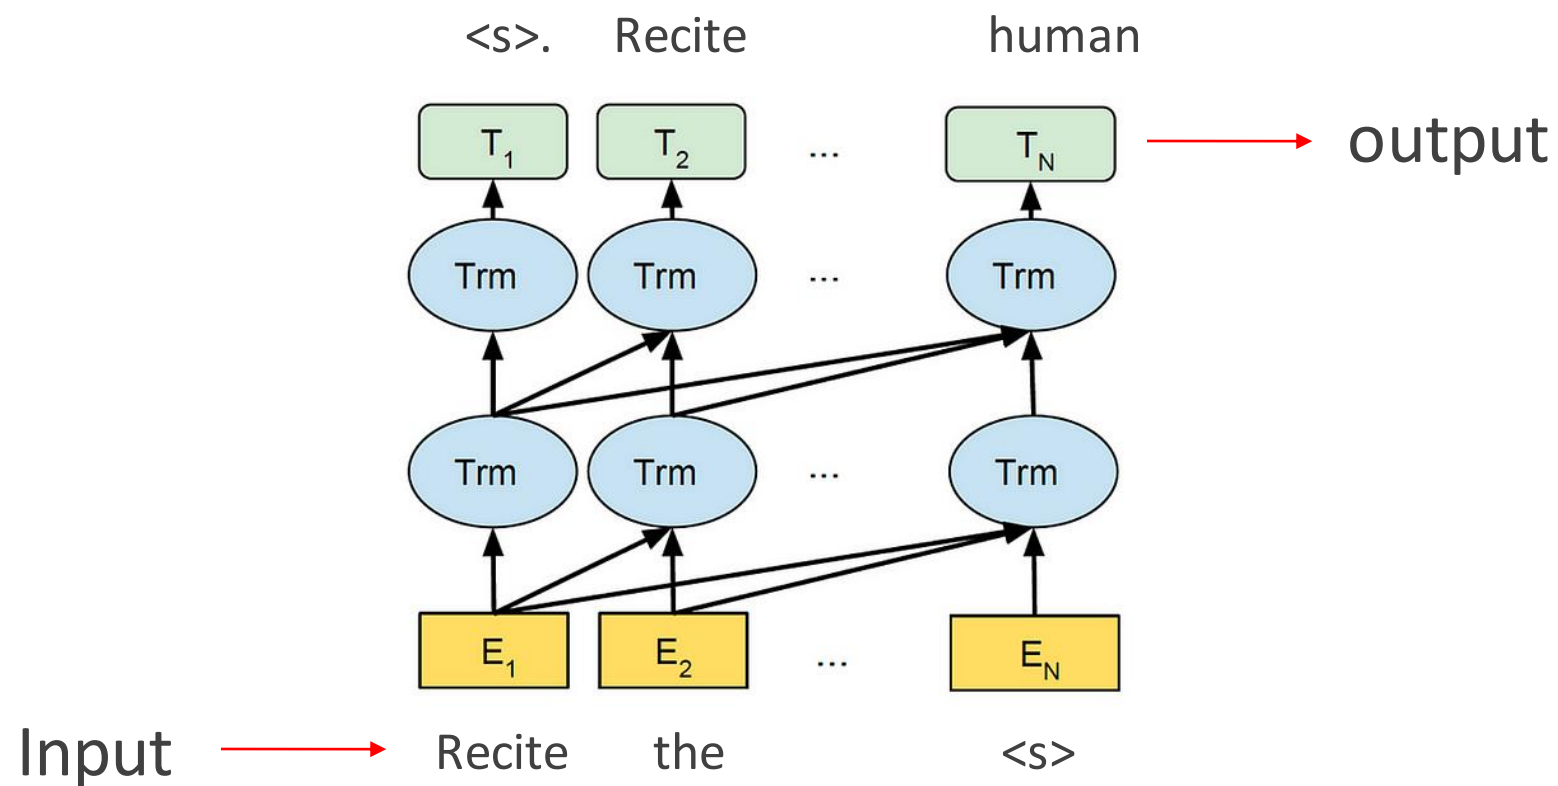
- Decoder-only
- The GPT-2 model consists of multiple layers of decoders.
- Masked Self-attention
 - The masking mechanism in GPT allows the model to look only forward.



GPT2 Generation



GPT2 Generation



Huggingface API

Install huggingface tools:

open the anaconda prompt

`activate [your_env_name]`

`pip install transformers`

`pip install datasets`

"transformers" provides tokenizer and models.

"datasets" provides datasets.

<https://huggingface.co/>

Tutorials:

<https://huggingface.co/docs>



Load Model&Tokenizer

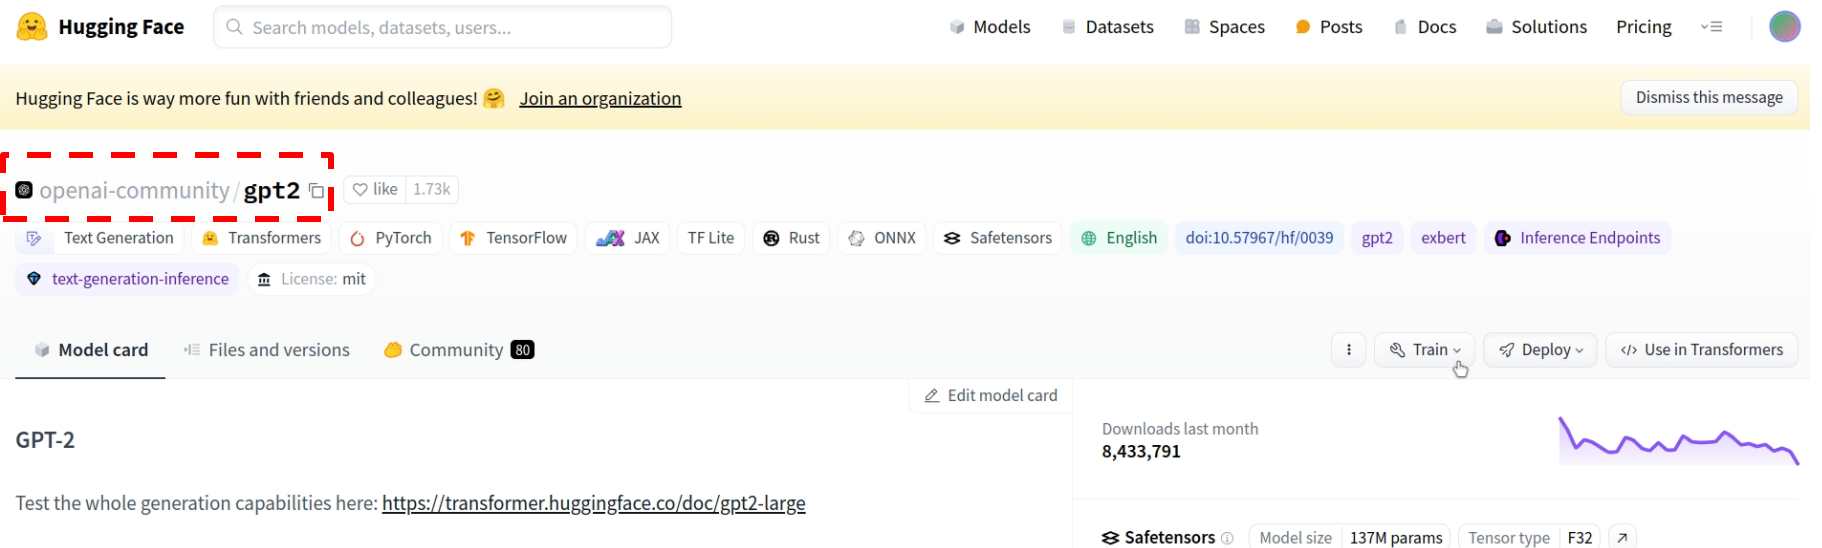
Load model:

```
gpt2_model = transformers.GPT2LMHeadModel.from_pretrained("model_name")
```

Load tokenizer:

```
gpt2_tokenizer = transformers.AutoTokenizer.from_pretrained("model_name")
```

model name is here



Hugging Face

Search models, datasets, users...

Models Datasets Spaces Posts Docs Solutions Pricing

Hugging Face is way more fun with friends and colleagues! [Join an organization](#)

Dismiss this message

openai-community / gpt2 like 1.73k

Text Generation Transformers PyTorch TensorFlow JAX TF Lite Rust ONNX Safetensors English doi:10.57967/hf/0039 gpt2 exbert Inference Endpoints

text-generation-inference License: mit

Model card Files and versions Community 80

GPT-2

Test the whole generation capabilities here: <https://transformer.huggingface.co/doc/gpt2-large>


Downloads last month
8,433,791

Safetensors Model size 137M params Tensor type F32



Load Model&Tokenizer


In text generation, sometimes we should set `padding_side='left'`, when loading tokenizer.



<s>	Recite	the	first	law	<s>	<pad>	<pad>	<pad>	
<s>	How	are	you	<s>	<pad>	<pad>	<pad>	<pad>	
<s>	Who	is	the	first	president	of	U.S.	<s>	

If padding is at the end the prompt, it would be unreasonable.

 **Left-padding**



<pad>	<pad>	<pad>	<s>	Recite	the	first	law	<s>	
<pad>	<pad>	<pad>	<pad>	<s>	How	are	you	<s>	
<s>	Who	is	the	first	president	of	U.S.	<s>	

Aligning the new tokens.

Tokenize the sequence

The tokenizer API can transform the text into ids.

It can also pack the text batch into tensors:

```
gpt2_tokenizer.batch_encode_plus(text, padding=True,  
                                truncation=True, return_tensors="pt")
```

`batch_encode_plus` returns a "BatchEncoding" object which can be regarded as dict.

(located on CPU ram)

Forward

Input the "BatchEncoding" object to the model

```
loss = gpt2_model(**inputs, labels=inputs["input_ids"]).loss
```

If the model is initialized for a specific task (e.g. `GPT2ForSequenceClassification`), the label can be directly input here and the API compute the loss automatically.

However, if you just:

```
gpt2_model = T.GPT2Model.from_pretrained("model_name")
```

Then you have to define the loss function by yourself.

e.g. `torch.nn.CrossEntropyLoss()`



Visualization

tqdm

During training, tqdm can provide a visual **progress bar**.

A simple example

import package

Process bar

```
from tqdm import tqdm
list = [1,2,3,4,5,6]
bar = tqdm(list)
for i in bar:
    pass
```

```
100%|██████████| 6/6 [00:00<00:00, 15069.36it/s]
```

tqdm

Set description to the front:

```
bar.set_description("List number: ")
```

```
List number: : 100%|██████████| 6/6 [00:00<00:00, 4707.41it/s]
```

Add item to the end:

```
bar.set_postfix(number=i)
```

```
List number: : 100%|██████████| 6/6 [00:00<00:00, 750.90it/s, number=6]
```

Tools for recording results

Tensorboard

Install:

```
pip install tensorboard
```

Import:

```
from torch.utils import tensorboard as tb
```

Documentation:

<https://pytorch.org/docs/stable/tensorboard.html>

Weights & Bias

Install:

```
pip install wandb
```

Import:

```
import wandb
```

Documentation:

<https://docs.wandb.ai/>

